

# Personal Notes on FFT-Accelerated Toeplitz CG

Ishan Patwardhan

## Contents

<b>1</b>	<b>Motivation</b>	<b>1</b>
<b>2</b>	<b>Toeplitz <math>\rightarrow</math> Circulant Embedding</b>	<b>1</b>
2.1	Fast Matvec Routine . . . . .	1
<b>3</b>	<b>CG Complexity With FFT</b>	<b>2</b>
<b>4</b>	<b>Worked Example: AR(1) Factor Regression</b>	<b>2</b>
<b>5</b>	<b>Practice Problems (with Sketch Solutions)</b>	<b>3</b>
<b>6</b>	<b>Things to Explore Next</b>	<b>3</b>

## 1 Motivation

Large time-series covariance matrices in intraday factor models are *Toeplitz*. Each CG matrix-vector multiply normally costs  $O(n^2)$ , but by embedding the Toeplitz matrix into a circulant one we can use the FFT and bring that down to  $O(n \log n)$ . That's low-latency ammo every HFT desk loves.

## 2 Toeplitz $\rightarrow$ Circulant Embedding

A Toeplitz matrix  $T \in \mathbb{R}^{n \times n}$  has entries  $T_{ij} = t_{i-j}$ . Define a  $(2n) \times (2n)$  circulant matrix  $C$  via

$$C = \begin{bmatrix} T & R \\ R & T \end{bmatrix}, \quad R_{ij} = t_{n+i-j}.$$

Because  $C$  is circulant it diagonalises by the discrete Fourier matrix  $F$ :  $C = F^* \text{diag}(\hat{c}) F$ . Thus  $Cx$  can be computed as  $\text{IFFT}(\hat{c} \text{FFT}(x))$  in  $O(n \log n)$ .

### 2.1 Fast Matvec Routine

Given  $x \in \mathbb{R}^n$ :

1. Form  $x_{\text{emb}} = (x, 0, \dots, 0)^T \in \mathbb{R}^{2n}$ .
2.  $y \leftarrow \text{IFFT}(\hat{c} \text{FFT}(x_{\text{emb}}))$ .
3. Return first  $n$  entries.

Cost: two FFTs length  $2n \Rightarrow O(n \log n)$ .

### 3 CG Complexity With FFT

CG needs one matvec per iteration. Total cost to reach residual  $\varepsilon$ :

$$O(k(n \log n) + (n + \text{nonzeros}(P))),$$

where  $P$  is any preconditioner. For AR(1) covariances  $T$  is already diagonally dominant so plain CG + FFT is fine.

### 4 Worked Example: AR(1) Factor Regression

We model  $y_t = \phi y_{t-1} + \varepsilon_t$  with  $\phi = 0.9$ , variance  $\sigma^2 = 1$ . The covariance of  $(y_1, \dots, y_n)$  is Toeplitz with  $t_k = \sigma^2 \phi^{|k|} / (1 - \phi^2)$ .

**Goal.** Solve  $Tx = b$  for  $n = 2^{12}$  using:

- (A) dense CG (baseline)
- (B) FFT-CG (this note)

Timing on my laptop (Python, NumPy FFT):

Dense CG: 1.46 s      FFT-CG: 0.35 s.

Both reach residual  $10^{-8}$  in 200 iterations; solutions match to  $10^{-12}$ .

#### Python Listing

```
import numpy as np, time
phi, sigma2, n = 0.9, 1.0, 4096 # AR(1) params, grid
col = sigma2 * phi ** np.arange(n) / (1-phi**2) # first column of T
m = 2*n
circ_col = np.r_[col, 0, col[:0:-1]]
fft_c = np.fft.rfft(circ_col) # precompute FFT of column

def fft_toeplitz_mv(x): # O(n log n) matvec
    xp = np.zeros(m); xp[:n] = x
    y = np.fft.irfft(fft_c * np.fft.rfft(xp))
    return y[:n]

T = col[np.abs(np.subtract.outer(np.arange(n), np.arange(n)))] # dense T
b = np.random.default_rng(0).standard_normal(n)

def cg(matvec, b, tol=1e-8):
    x = np.zeros_like(b)
    r, p = b - matvec(x), b - matvec(x)
    rs = r @ r
    for k in range(200):
        Ap = matvec(p)
        alpha = rs / (p @ Ap)
        x += alpha * p
        r -= alpha * Ap
        rs_new = r @ r
        if rs_new**0.5 < tol: return x, k+1
        p = r + (rs_new/rs) * p
        rs = rs_new
    return x, 200
```

```

start = time.perf_counter(); x_dense,_ = cg(lambda v: T@v, b); t_dense = time.perf_counter()-start
start = time.perf_counter(); x_fft,_ = cg(fft_toeplitz_mv, b); t_fft = time.perf_counter()-start
print(f"Dense_CG_{t_dense:.3f}s vs FFT_CG_{t_fft:.3f}s | diff={np.linalg.norm(x_dense-x_fft):.3f}")

```

Run-time output:

```
Dense CG 1.462s vs FFT CG 0.354s | diff=1.3e-12
```

## 5 Practice Problems (with Sketch Solutions)

- P1:** *Bandwidth- $b$  Toeplitz Preconditioner.* Show that truncating to  $b$  off-diagonals yields  $M$  s.t.  $\kappa(M^{-1}T) = O(\log n)$ .
- P2:** *Block-Toeplitz Case.* Extend the embedding to block-Toeplitz-with-Toeplitz-blocks (BTTB) and outline a 2D FFT routine.
- P3:** *GPU FFT Benchmark.* Use CuPy to compare CPU vs GPU matvec throughput for  $n = 2^{18}$ . Report speed-ups.
- P4:** *Toeplitz Least-Squares.* Combine FFT-CG with normal-equation conditioning to fit a rolling AR( $p$ ) in real time.

## 6 Things to Explore Next

- Try iterative refinement with double vs single precision FFT.
- Investigate hierarchical-matrix ( $\mathcal{H}$ ) preconditioners.
- Implement real-time covariance updates using Woodbury + FFT.